

```

import sys
import numpy as np
import pandas as pd
import seaborn as sns
import pickle
%matplotlib inline
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
import sklearn.metrics as metrics

df=pd.read_csv('C:\\Users\\ketziyal\\Downloads\\flightdata.csv')

```

#Analyze the data

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11231 entries, 0 to 11230
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   YEAR                                11231 non-null  int64
1   QUARTER                            11231 non-null  int64
2   MONTH                              11231 non-null  int64
3   DAY_OF_MONTH                       11231 non-null  int64
4   DAY_OF_WEEK                        11231 non-null  int64
5   UNIQUE_CARRIER                   11231 non-null  object
6   TAIL_NUM                           11231 non-null  object
7   FL_NUM                             11231 non-null  int64
8   ORIGIN_AIRPORT_ID                 11231 non-null  int64
9   ORIGIN                             11231 non-null  object
10  DEST_AIRPORT_ID                   11231 non-null  int64
11  DEST                              11231 non-null  object
12  CRS_DEP_TIME                      11231 non-null  int64
13  DEP_TIME                          11124 non-null  float64
14  DEP_DELAY                         11124 non-null  float64
15  DEP_DEL15                         11124 non-null  float64
16  CRS_ARR_TIME                      11231 non-null  int64
17  ARR_TIME                          11116 non-null  float64
18  ARR_DELAY                         11043 non-null  float64
19  ARR_DEL15                         11043 non-null  float64
20  CANCELLED                         11231 non-null  float64
21  DIVERTED                          11231 non-null  float64
22  CRS_ELAPSED_TIME                  11231 non-null  float64
23  ACTUAL_ELAPSED_TIME              11043 non-null  float64
24  DISTANCE                          11231 non-null  float64
25  Unnamed: 25                      0 non-null     float64

```

```
dtypes: float64(12), int64(10), object(4)
memory usage: 2.2+ MB
```

```
df.describe()
```

	YEAR	QUARTER	MONTH	DAY_OF_MONTH	DAY_OF_WEEK
\count	11231.0	11231.000000	11231.000000	11231.000000	11231.000000
mean	2016.0	2.544475	6.628973	15.790758	3.960199
std	0.0	1.090701	3.354678	8.782056	1.995257
min	2016.0	1.000000	1.000000	1.000000	1.000000
25%	2016.0	2.000000	4.000000	8.000000	2.000000
50%	2016.0	3.000000	7.000000	16.000000	4.000000
75%	2016.0	3.000000	9.000000	23.000000	6.000000
max	2016.0	4.000000	12.000000	31.000000	7.000000

	FL_NUM	ORIGIN_AIRPORT_ID	DEST_AIRPORT_ID	CRS_DEP_TIME
\count	11231.000000	11231.000000	11231.000000	11231.000000
mean	1334.325617	12334.516695	12302.274508	1320.798326
std	811.875227	1595.026510	1601.988550	490.737845
min	7.000000	10397.000000	10397.000000	10.000000
25%	624.000000	10397.000000	10397.000000	905.000000
50%	1267.000000	12478.000000	12478.000000	1320.000000
75%	2032.000000	13487.000000	13487.000000	1735.000000
max	2853.000000	14747.000000	14747.000000	2359.000000

	DEP_TIME	...	CRS_ARR_TIME	ARR_TIME	ARR_DELAY	\
count	11124.000000	...	11231.000000	11116.000000	11043.000000	
mean	1327.189410	...	1537.312795	1523.978499	-2.573123	
std	500.306462	...	502.512494	512.536041	39.232521	
min	1.000000	...	2.000000	1.000000	-67.000000	
25%	905.000000	...	1130.000000	1135.000000	-19.000000	

50%	1324.000000	...	1559.000000	1547.000000	-10.000000
75%	1739.000000	...	1952.000000	1945.000000	1.000000
max	2400.000000	...	2359.000000	2400.000000	615.000000

	ARR_DEL15	CANCELLED	DIVERTED	CRS_ELAPSED_TIME \
count	11043.000000	11231.000000	11231.000000	11231.000000
mean	0.124513	0.010150	0.006589	190.652124
std	0.330181	0.100241	0.080908	78.386317
min	0.000000	0.000000	0.000000	93.000000
25%	0.000000	0.000000	0.000000	127.000000
50%	0.000000	0.000000	0.000000	159.000000
75%	0.000000	0.000000	0.000000	255.000000
max	1.000000	1.000000	1.000000	397.000000

	ACTUAL_ELAPSED_TIME	DISTANCE	Unnamed: 25
count	11043.000000	11231.000000	0.0
mean	179.661233	1161.031965	NaN
std	77.940399	643.683379	NaN
min	75.000000	509.000000	NaN
25%	117.000000	594.000000	NaN
50%	149.000000	907.000000	NaN
75%	236.000000	1927.000000	NaN
max	428.000000	2422.000000	NaN

[8 rows x 22 columns]

#handling missing values

df.isnull().sum()

YEAR	0
QUARTER	0
MONTH	0
DAY_OF_MONTH	0
DAY_OF_WEEK	0
UNIQUE_CARRIER	0
TAIL_NUM	0
FL_NUM	0
ORIGIN_AIRPORT_ID	0
ORIGIN	0
DEST_AIRPORT_ID	0
DEST	0
CRS_DEP_TIME	0
DEP_TIME	107
DEP_DELAY	107
DEP_DEL15	107
CRS_ARR_TIME	0
ARR_TIME	115
ARR_DELAY	188
ARR_DEL15	188

```
CANCELLED          0
DIVERTED           0
CRS_ELAPSED_TIME   0
ACTUAL_ELAPSED_TIME 188
DISTANCE           0
Unnamed: 25        11231
dtype: int64
```

```
df['DEST'].unique()
```

```
array(['SEA', 'MSP', 'DTW', 'ATL', 'JFK'], dtype=object)
```

```
#data visualization
```

```
from matplotlib import pyplot as plt
```

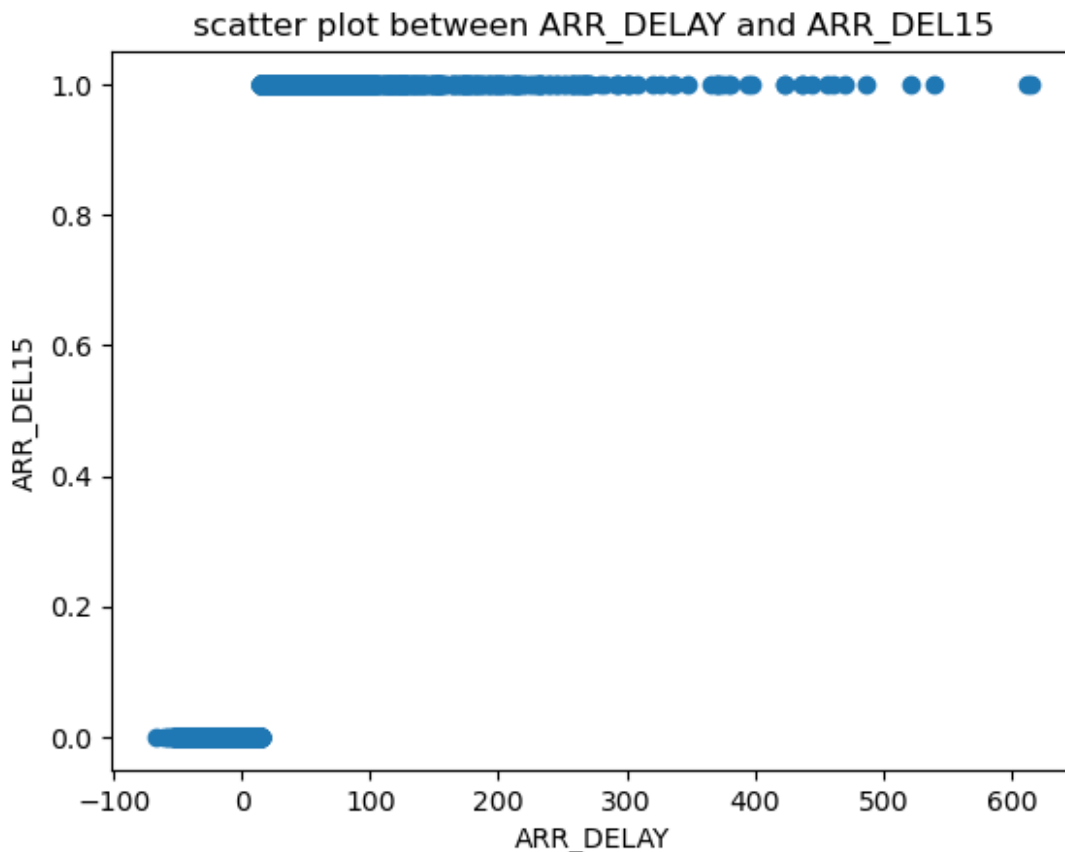
```
plt.scatter(df['ARR_DELAY'],df['ARR_DEL15'])
```

```
plt.xlabel('ARR_DELAY')
```

```
plt.ylabel('ARR_DEL15')
```


```
plt.title('scatter plot between ARR_DELAY and ARR_DEL15')
```

```
Text(0.5, 1.0, 'scatter plot between ARR_DELAY and ARR_DEL15')
```




```
sns.catplot(x="ARR_DEL15",y="ARR_DELAY",kind='bar',data=df)
```

<seaborn.axisgrid.FacetGrid at 0x196d4d7be50>



```
sns.heatmap(df.corr())
```

<AxesSubplot:>



```
df=df.drop('Unnamed: 25',axis=1)
df.isnull().sum()
```

```
YEAR          0
QUARTER       0
MONTH         0
DAY_OF_MONTH  0
DAY_OF_WEEK   0
UNIQUE_CARRIER  0
TAIL_NUM      0
FL_NUM        0
ORIGIN_AIRPORT_ID  0
ORIGIN         0
DEST_AIRPORT_ID  0
DEST           0
CRS_DEP_TIME   0
DEP_TIME       107
DEP_DELAY      107
DEP_DEL15      107
CRS_ARR_TIME   0
ARR_TIME       115
ARR_DELAY      188
```

```

ARR_DEL15          188
CANCELLED          0
DIVERTED           0
CRS_ELAPSED_TIME   0
ACTUAL_ELAPSED_TIME 188
DISTANCE           0
dtype: int64

```

```

df=df[["FL_NUM","MONTH","DAY_OF_MONTH","DAY_OF_WEEK","ORIGIN","DEST","
CRS_ARR_TIME","DEP_DEL15","ARR_DEL15"]]
df.isnull().sum()

```

```

FL_NUM          0
MONTH           0
DAY_OF_MONTH    0
DAY_OF_WEEK     0
ORIGIN          0
DEST            0
CRS_ARR_TIME    0
DEP_DEL15      107
ARR_DEL15      188
dtype: int64

```

```

df=df.fillna({'ARR_DEL15':1})
df=df.fillna({'DEP_DEL15':0})
df.iloc[177:185]

```

	FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	ORIGIN	DEST
CRS_ARR_TIME \						
177	2834	1	9	6	MSP	SEA
852						
178	2839	1	9	6	DTW	JFK
1724						
179	86	1	10	7	MSP	DTW
1632						
180	87	1	10	7	DTW	MSP
1649						
181	423	1	10	7	JFK	ATL
1600						
182	440	1	10	7	JFK	ATL
849						
183	485	1	10	7	JFK	SEA
1945						
184	557	1	10	7	MSP	DTW
912						

	DEP_DEL15	ARR_DEL15
177	0.0	1.0
178	0.0	0.0
179	0.0	1.0
180	1.0	0.0

```

181         0.0         0.0
182         0.0         0.0
183         1.0         0.0
184         0.0         1.0

```

```
import math
```

```

for index,row in df.iterrows():
    df.loc[index,'CRS_ARR_TIME']=math.floor(row['CRS_ARR_TIME']/100)
df.head()

```

```

      FL_NUM  MONTH  DAY_OF_MONTH  DAY_OF_WEEK  ORIGIN  DEST  CRS_ARR_TIME
\
0    1399      1          1          5    ATL  SEA          21
1    1476      1          1          5    DTW  MSP          14
2    1597      1          1          5    ATL  SEA          12
3    1768      1          1          5    SEA  MSP          13
4    1823      1          1          5    SEA  DTW           6

```

```

      DEP_DEL15  ARR_DEL15
0          0.0          0.0
1          0.0          0.0
2          0.0          0.0
3          0.0          0.0
4          0.0          0.0

```

```

from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()

```

```

df['DEST']=le.fit_transform(df['DEST'])
df['ORIGIN']=le.fit_transform(df['ORIGIN'])

```

```
df.head()
```

```

      FL_NUM  MONTH  DAY_OF_MONTH  DAY_OF_WEEK  ORIGIN  DEST
CRS_ARR_TIME \
0    1399      1          1          5          0      4
21
1    1476      1          1          5          1      3
14
2    1597      1          1          5          0      4
12
3    1768      1          1          5          4      3
13
4    1823      1          1          5          4      1
6

```


	DEP_DEL15	ARR_DEL15
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

```
from sklearn.preprocessing import OneHotEncoder
```

```
oh=OneHotEncoder()
```

```
z=oh.fit_transform(x[:,4:5]).toarray()
```

```
t=oh.fit_transform(x[:,5:6]).toarray()
```

```
z
```

```
array([[1., 0., 0., 0., 0.],
       [0., 1., 0., 0., 0.],
       [1., 0., 0., 0., 0.],
       ...,
       [0., 1., 0., 0., 0.],
       [1., 0., 0., 0., 0.],
       [1., 0., 0., 0., 0.]])
```

```
t
```

```
array([[0., 0., 0., 0., 1.],
       [0., 0., 0., 1., 0.],
       [0., 0., 0., 0., 1.],
       ...,
       [0., 0., 0., 0., 1.],
       [0., 0., 0., 0., 1.],
       [0., 1., 0., 0., 0.]])
```

```
df=pd.get_dummies(df,columns=['ORIGIN','DEST'])
```

```
df.head()
```

	FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	CRS_ARR_TIME	
DEP_DEL15						
0	1399	1	1	5	21	0.0
1	1476	1	1	5	14	0.0
2	1597	1	1	5	12	0.0
3	1768	1	1	5	13	0.0
4	1823	1	1	5	6	0.0

	ARR_DEL15	ORIGIN_0	ORIGIN_1	ORIGIN_2	ORIGIN_3	ORIGIN_4	DEST_0
0	0.0	1	0	0	0	0	0
1	0.0	0	1	0	0	0	0
2	0.0	1	0	0	0	0	0
3	0.0	0	0	0	0	1	0
4	0.0	0	0	0	0	1	0

	DEST_1	DEST_2	DEST_3	DEST_4
0	0	0	0	1
1	0	0	1	0
2	0	0	0	1
3	0	0	1	0
4	1	0	0	0

```
x=df.iloc[:,0:8].values
y=df.iloc[:,8:9].values
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

```
x_test.shape
```

```
(2247, 8)
```

```
x_train.shape
```

```
(8984, 8)
```

```
y_test.shape
```

```
(2247, 1)
```

```
y_train.shape
```

```
(8984, 1)
```

MODEL BUILDING

```
#decision tree
```

```
from sklearn.tree import DecisionTreeClassifier
dc=DecisionTreeClassifier()
dc.fit(x_train,y_train)
dc.score(x_test,y_test)
```

0.8607031597685804

#random forest

```
from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(n_estimators=50,random_state=42)
rf.fit(x_train,y_train)
rf.score(x_test,y_test)
```

C:\Users\ketziyal\AppData\Local\Temp\ipykernel_19504\905497165.py:3:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples,), for
example using ravel().

```
    rf.fit(x_train,y_train)
```

0.910547396528705

```
pd.DataFrame(rf.predict(x_test)).value_counts()
```

```
0.0    2003
1.0     244
dtype: int64
```

#logestic regression

```
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression(solver='sag')
lr.fit(x_train,y_train)
lr.score(x_test,y_test)
```

c:\Users\ketziyal\anaconda3\lib\site-packages\sklearn\utils\
validation.py:1111: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples,), for example using ravel().

```
    y = column_or_1d(y, warn=True)
```

c:\Users\ketziyal\anaconda3\lib\site-packages\sklearn\linear_model\
_sag.py:350: ConvergenceWarning: The max_iter was reached which means
the coef_ did not converge
 warnings.warn(

0.8615932354250111

```
lr.predict(x_test).sum()
```

0.0

#svm

```
from sklearn.svm import SVC
svm=SVC(kernel='sigmoid')
svm.fit(x_train,y_train)
svm.score(x_test,y_test)
```

```
c:\Users\ketziyal\anaconda3\lib\site-packages\sklearn\utils\
validation.py:1111: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
```

```
0.7725856697819314
```

```
pd.DataFrame(svm.predict(x_test)).value_counts()
```

```
0.0    1941
1.0     306
dtype: int64
```

```
pd.DataFrame(y_test).value_counts()
```

```
0.0    1936
1.0     311
dtype: int64
```

```
# K-NEAREST NEIGHBOUR CLASSIFIER
```

```
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=5)
knn.fit(x_train,y_train)
knn.score(x_test,y_test)
```

```
c:\Users\ketziyal\anaconda3\lib\site-packages\sklearn\neighbors\
_classification.py:207: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    return self._fit(X, y)
```

```
0.8486871384067646
```

```
pd.DataFrame(knn.predict(x_test)).value_counts()
```

Evaluation Of Random Forest


```
from sklearn.metrics import
confusion_matrix,accuracy_score,classification_report
pred=rf.predict(x_test)
cm=confusion_matrix(y_test, pred)
plt.figure(figsize=(10,6))
sns.heatmap(cm, annot=True,cmap='winter',linewidths=0.3,
linecolor='black',annot_kws={"size": 20})
TP=cm[0][0]
TN=cm[1][1]
FN=cm[1][0]
FP=cm[0][1]
#print(round(accuracy_score(prediction3,y_test)*100,2))
#print('Testing Accuracy for knn',(TP+TN)/(TP+TN+FN+FP))
```

```

print('Testing Sensitivity for Random Forest',(TP/(TP+FN)))
print('Testing Specificity for Random Forest',(TN/(TN+FP)))
print('Testing Precision for Random Forest',(TP/(TP+FP)))
print('Testing accuracy for Random Forest',accuracy_score(y_test,
pred))

```

Testing Sensitivity for Random Forest 0.9331003494757864
 Testing Specificity for Random Forest 0.7254098360655737
 Testing Precision for Random Forest 0.9653925619834711
 Testing accuracy for Random Forest 0.910547396528705



```

print(classification_report(y_test,pred))

```

	precision	recall	f1-score	support
0.0	0.93	0.97	0.95	1936
1.0	0.73	0.57	0.64	311
accuracy			0.91	2247
macro avg	0.83	0.77	0.79	2247
weighted avg	0.90	0.91	0.91	2247

#Evaluation Of Decission Tree

```

pred1=dc.predict(x_test)
cm1=confusion_matrix(y_test, pred1)
plt.figure(figsize=(10,6))
sns.heatmap(cm1, annot=True,cmap='winter',linewidths=0.3,

```

```

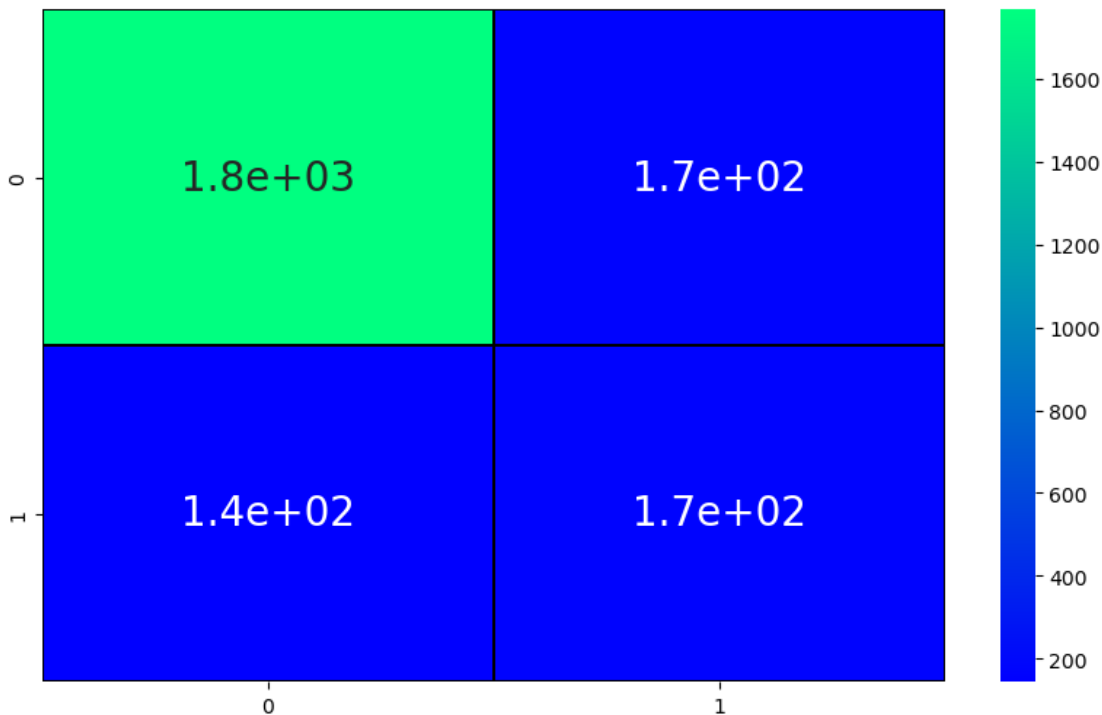
linecolor='black',annot_kws={"size": 20})
TP=cm1[0][0]
TN=cm1[1][1]
FN=cm1[1][0]
FP=cm1[0][1]
#print(round(accuracy_score(prediction3,y_test)*100,2))
print('Testing Accuracy for Decision Tree',(TP+TN)/(TP+TN+FN+FP))
print('Testing Sensitivity for Decision Tree',(TP/(TP+FN)))
print('Testing Specificity for Decision Tree',(TN/(TN+FP)))
print('Testing Precision for Decision Tree',(TP/(TP+FP)))
print('Testing accuracy for Decision Tree',accuracy_score(y_test,
pred1))

```

```

Testing Accuracy for Decision Tree 0.8607031597685804
Testing Sensitivity for Decision Tree 0.9246467817896389
Testing Specificity for Decision Tree 0.49702380952380953
Testing Precision for Decision Tree 0.9127066115702479
Testing accuracy for Decision Tree 0.8607031597685804

```



```
print(classification_report(y_test,pred1))
```

	precision	recall	f1-score	support
0.0	0.92	0.91	0.92	1936
1.0	0.50	0.54	0.52	311
accuracy			0.86	2247
macro avg	0.71	0.72	0.72	2247

weighted avg	0.87	0.86	0.86	2247
--------------	------	------	------	------

```
import pickle
```

```
pickle.dump(rf,open("flight.pkl",'wb'))
```